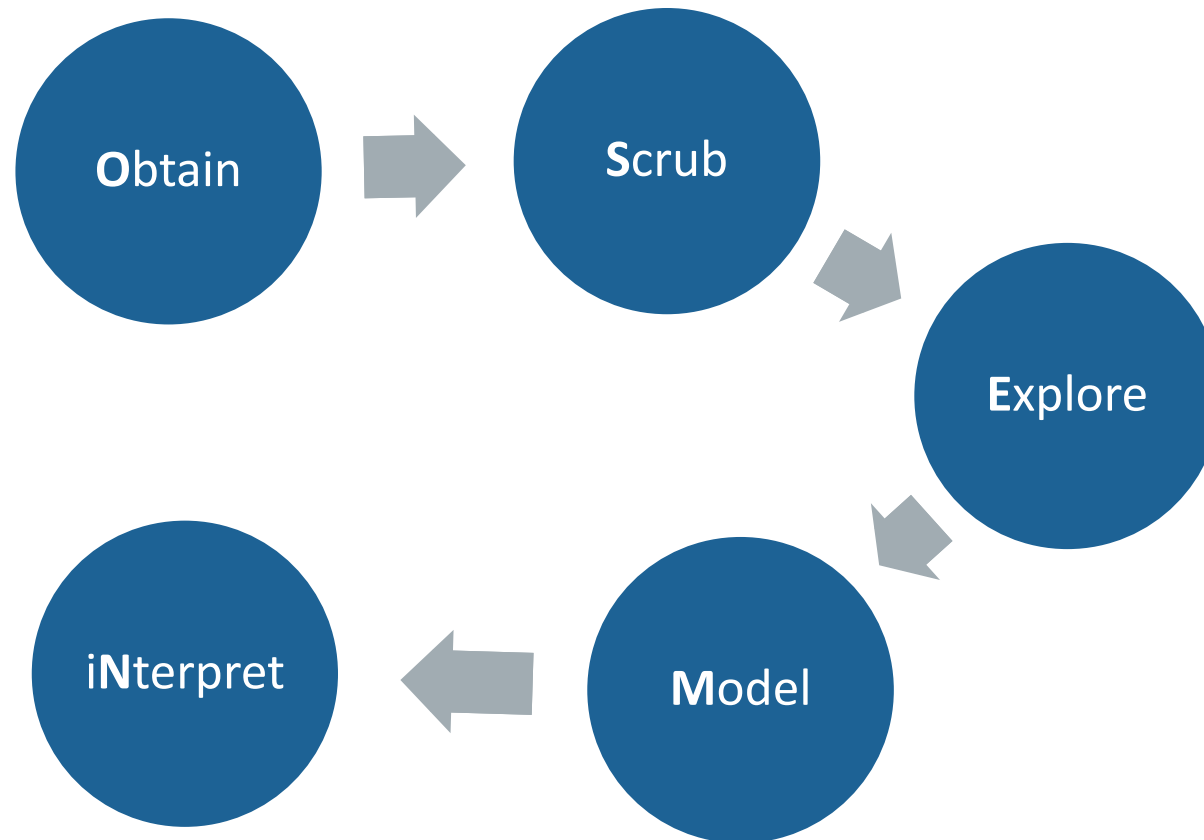




Machine Learning Overview Part I

School of Information Studies
Syracuse University


The Machine Learning Process OSEMN (Rhymes With Possum)



OSEMN Phase: Obtain

Collect data


- Deal with file formats and how to read data
- Query databases or data repositories
- Extract data from other sources
- Generate data (e.g., surveys, sensors)



OSEMN Phase: Scrub

Get data into a useable format structure


- Filter/subset
- Extract attributes
- Replace/handle missing, illegal, and anomalous values
- Transform/bin/code data attributes



OSEMN Phase: Explore

Explore patterns and trends

- Start to “understand” the data, detect outliers
- Visualize attributes (e.g., scatter plots, histograms)
- Calculate/visualize descriptive statistics (e.g., distributions)
- Feature selection (what attributes are most interesting)

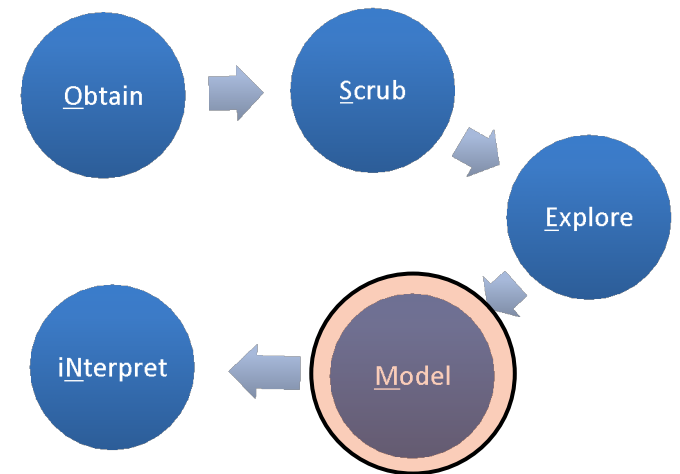


OSEMN Phase: Model

Build predictive models (machine learning)

- Type of modeling:
 - Supervised (classification, regression)
 - Unsupervised (e.g., clustering)


→ *Will discuss shortly*
- Model tuning and comparing candidate models



OSEMN Phase: Interpret

Understand/explain the results

- Draw conclusions from and data models
- Evaluate the meaning of results
- Communicate the results
- Ensure actionable insight



Question

How is the OSEMN process different from or the same as the overall data science process?




Machine Learning Overview

Part II

School of Information Studies
Syracuse University

Question

How is the OSEMN process different from or the same as the overall data science process?




Supervised Machine Learning in a Nutshell

“Supervised” refers to the idea that there is a **criterion** used during an algorithm training phase.

A supervised algorithm uses a set of input variables to optimize the prediction of an outcome variable.

Supervised Learning



Unsupervised Learning


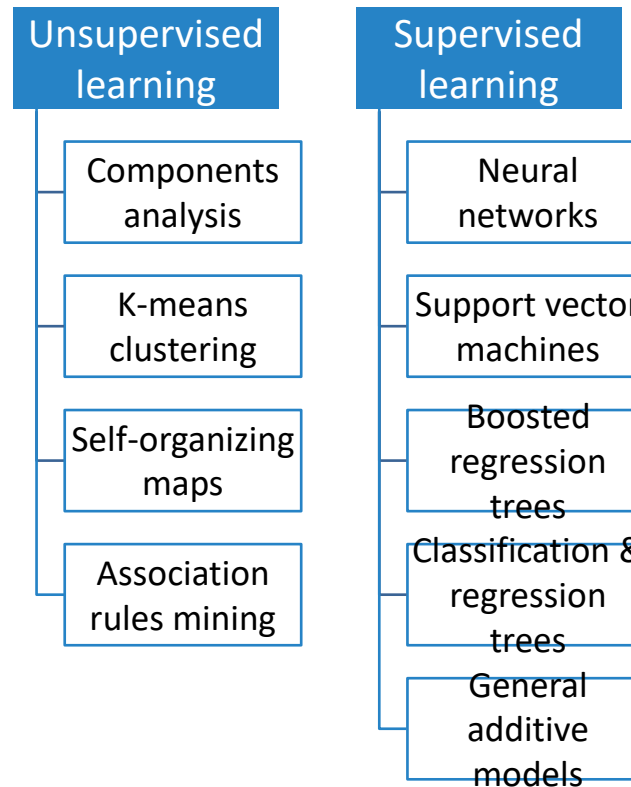


Image credit: Olivia Klose

Examples of Machine Learning Techniques

Unsupervised learning includes a variety of machine learning techniques that do not use a criterion or dependent variable, but rather look for patterns solely among “independent” variables.



Supervised learning is parallel in concept to the predictive statistical techniques used by many social science researchers, such as linear regression, but without the restriction of only exploring linear relationships.

Another form of learning is known as “**reinforcement learning**.” Evolution of these models depends on success/failure cues from real or simulated environments.

Question

What is the meaning of the word “supervised” in the context of data mining?



Machine Learning Overview

Part III

School of Information Studies
Syracuse University

Two Types of Prediction (For Supervised Learning)

1. Classification problems

- **Predicting membership in two or more categories** based on a set of predictors (model features)
 - **Examples of criteria to predict**
 - Medical diagnosis
 - Employment outcomes (e.g., hiring)
 - Financial outcome (e.g., customer default on a loan)

2. Regression problems

- **Predicting a continuous numeric outcome** based on a set of predictors. **Examples of criteria to predict:** sales volume, employee engagement, customer satisfaction

Traditional approaches

Classification problems

- Logistic regression
- Discriminant analysis

Regression problems

- Ordinary least squares regression
(lm() models)
- Generalized linear models
- Lasso regression

Supervised Learning Example

Train a machine learning algorithm to predict the weather

- Collect weather data over a period of time
 - Sunny, cloudy
 - Temperature
 - Barometer
 - Wind speed and direction
- Train a machine learning algorithm with these collected variables
- Collect more weather data and predict the weather via our trained algorithm
 - Classification would be predicting good weather or bad weather
 - Regression would be predicting the temperature
- Then validate the prediction

The Modeling Process

Use a substantial number of training cases

- The machine learning algorithm can use that data to build a model

Use the results of this process (i.e., the model) on test data set to determine how well algorithm performed

- Validate the model on new data

The result is a model that can be used for prediction

- Predict data that was not used during training
- Predict future instances of data

*Note: The model is **not always useful** for explaining results to managers.*

In some cases, algorithms produce results that are not easy to interpret or visualize; for some algorithms there is no output that is like a regression coefficient.



Using Classification and Regression Trees

School of Information Studies
Syracuse University

An Example: CART

Classification and Regression Trees

A family of data mining techniques for developing predictions on **either** continuous outcome variables or class outcome variables

Can be used either for classification with unordered categories, or pseudo-continuous ordered outcomes

Uses iterative model building techniques, typically:

- Develops “splits” in the data where the level of a predictor is used to divide the data set into two (or more) partitions according to the status of the outcome variable
- The resulting models can be represented as binary trees

Example of rpart: For Titanic, Part I

```
library(rpart)
```

```
load("titanic.raw.rdata")
```

```
titanic <- titanic.raw
```

```
#Build the model
```

```
cartTree <- rpart(Survived ~ ., data = titanic)
```


Example of rpart: For Titanic, Part II

cartTree

n= 2201

node), split, n, loss, yval, (yprob)

*Denotes terminal node

- 1) root 2201 711 No (0.6769650 0.3230350)
- 2) Sex=Male 1731 367 No (0.7879838 0.2120162)
 - 4) Age=Adult 1667 338 No (0.7972406 0.2027594) *
 - 5) Age=Child 64 29 No (0.5468750 0.4531250)
 - 10) Class=3rd 48 13 No (0.7291667 0.2708333) *
 - 11) Class=1st,2nd 16 0 Yes (0.0000000 1.0000000) *
- 3) Sex=Female 470 126 Yes (0.2680851 0.7319149)
 - 6) Class=3rd 196 90 No (0.5408163 0.4591837) *
 - 7) Class=1st,2nd,Crew 274 20 Yes (0.0729927 0.9270073) *

Example of rpart: For Titanic, Part III

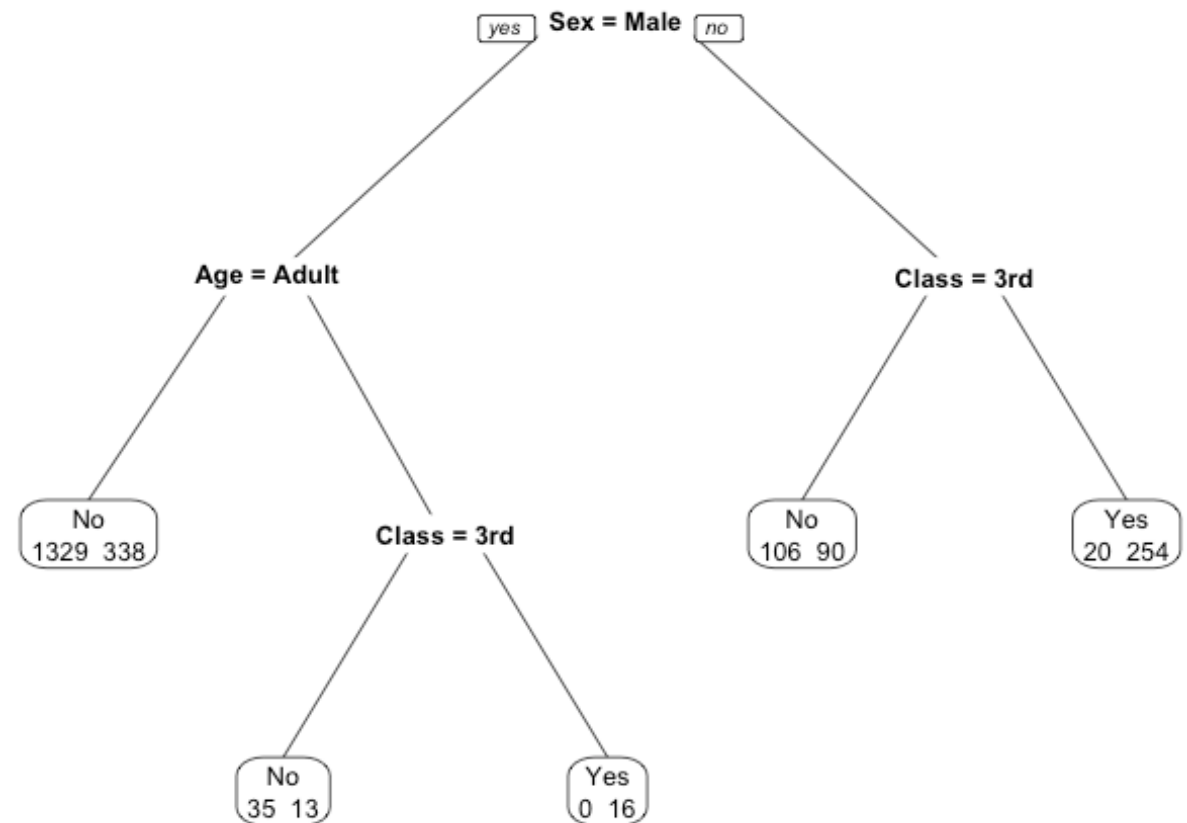
```
library(rpart.plot)
```

```
prp(cartTree, faclen = 0, cex = 0.8, extra = 1)
```

For left pointing leaves:

- The first number: number of correctly classified cases
- The second number: number of incorrectly classified cases

For right pointing leaves, vice versa.



Predicting Values Using the Model

```
predictedSurvived <- predict(cartTree, newdata=titanic, type = "class")
```

```
titanic[1,]
```

	Class	Sex	Age	Survived
3	3rd	Male	Child	No

```
predictedSurvived[1]
```

1

No

Levels: No Yes

What Was the Accuracy?

```
actualSurvived <- as.factor(titanic$Survived == "Yes")  
confMatrix <- table(predictedSurvived, actualSurvived)  
confMatrix
```

		actualSurvived	
predictedSurvived		FALSE	TRUE
No		1470	441
Yes		20	270

```
accuracy <- 1 - (sum(confMatrix) - sum(diag(confMatrix))) / sum(confMatrix)  
accuracy  
0.7905498
```

Questions

Is 79% good?



Using Classification and Regression Trees (cont.)

School of Information Studies
Syracuse University

Using Caret to Create Training and Test Sets

```
library(caret)
```

```
#Makes the sampling predictable
```

```
set.seed(111)
```

```
#Randomly sample elements to go into a training data set
```

```
trainList <- createDataPartition(y=titanic$Survived, p=.80, list=FALSE)
```

```
#Include all of those elements in the training set
```

```
trainSet <- titanic[trainList,]
```

```
#Construct test set from everything that didn't go into the training
```

```
testSet <- titanic[-trainList,]
```

Check Trainset

```
head(trainList)
```

```
  Resample1
```

```
[1,]      1
```

```
[2,]      2
```

```
[3,]      4
```

```
[4,]      5
```

```
[5,]      6
```


```
[6,]      7
```

Note that some
observations are
skipped.

Use the Trainset to Build a Model Using 'train'

```
cartTree <- train(Survived ~ ., data=trainSet,  
                  method="rpart")
```

```
prp(cartTree$finalModel, faclen = 0,  
    cex = 0.8, extra = 1)
```



Evaluate Using the Testset

#Note use of “raw”

```
predictValues <- predict(cartTree,newdata=testSet, type = "raw")
```

#Simpler to do confusion matrix

```
confusion <- confusionMatrix(predictValues, testSet$Survived)
```

```
confusion$overall[1]
```

Accuracy

0.7954545

Using all the data: the accuracy was
0.7905498

Exploring the Model Build via Caret

cartTree

CART

1761 samples

3 predictor

2 classes: 'No', 'Yes'

No pre-processing

Resampling: bootstrapped (25 reps)

Summary of sample sizes: 1761, 1761, 1761,...

Resampling results across tuning parameters:

...

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was $cp = 0.009666081$.

cartTree\$finalModel

n= 1761

1) root 1761 569 No (0.67688813 0.32311187)

2) SexMale \geq 0.5 1385 295 No (0.78700361 0.21299639) *

3) SexMale $<$ 0.5 376 102 Yes (0.27127660 0.72872340)

6) Class3rd \geq 0.5 156 73 No (0.53205128 0.46794872) *

7) Class3rd $<$ 0.5 220 19 Yes (0.08636364 0.91363636) *

Exploring the Model Build via Caret (cont.)

```
predictValues <- predict(cartTree1,newdata=testSet, type = "raw")
```

```
confusionMatrix(predictValues, testSet$Survived)
```

Confusion matrix and statistics

Prediction	Reference	
	FALSE	TRUE
FALSE	297	89
TRUE	1	53

Accuracy : 0.7955

95% CI : (0.7547, 0.8322)

No information rate : 0.6773

p value [Acc > NIR] : 2.382e-08

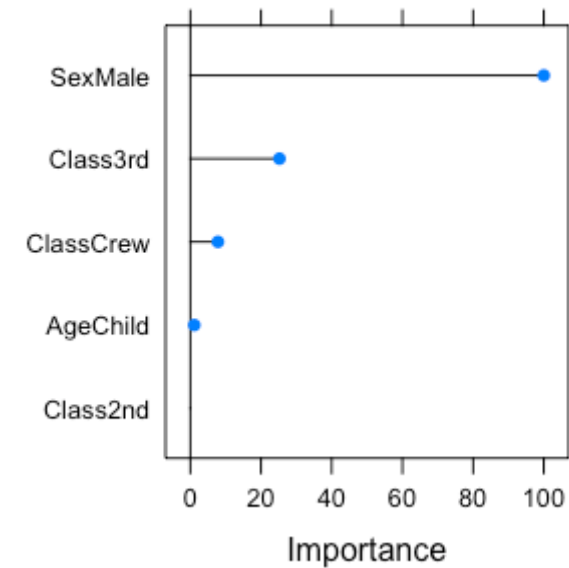
Exploring Variable Importance

```
varImp(cartTree)
```

rpart variable importance

	Overall
SexMale	100.000
Class3rd	25.233
ClassCrew	7.825
AgeChild	1.136
Class2nd	

```
plot(varImp(cartTree))
```





Advanced Tree Example

School of Information Studies
Syracuse University

The German Credit Database

From the UCI machine learning repository

Built-in data set stored in the caret package

1,000 cases and 62 variables

“Class” is the outcome variable: binary creditworthiness (“good” or “bad”)

Duration	Age	NumberExistingCredits	NumberPeopleMaintenance	Telephone	ForeignWorker	Class
4	67	2	1	0	1	Good
2	22	1	1	1	1	Bad
3	49	1	2	1	1	Good
4	45	1	2	1	1	Good
4	53	2	2	1	1	Bad
4	35	1	2	0	1	Good
4	53	1	1	1	1	Good
2	35	1	1	0	1	Good

Showing 1 to 9 of 1,000 entries

The German Credit Database (cont.)

Example predictors

- Checking account status, duration, credit history, purpose of the loan, amount of the loan, savings accounts or bonds, installment rate in percentage of disposable income, other installment plans, number of existing credits
- Employment duration, personal information, other debtors/guarantors, residence duration, property, age, housing, job information, number of people being liable to provide maintenance for, telephone, and foreign worker status

For convenience, we will only use the first 10 columns, including “Class”

We will try “treebag”, a bagged CART; there are more than 200 fitting algorithms

See <http://topepo.github.io/caret/modelList.html>

Create Training and Test Sets

```
#Just grab a subset of the data for the demo
```

```
data("GermanCredit")
```

```
subCredit <- GermanCredit[,1:10]
```

```
#Makes the sampling predictable
```

```
set.seed(111)
```

```
#Randomly sample elements to go into a training data set
```

```
trainList <- createDataPartition(y=subCredit$Class,p=.80,list=FALSE)
```

```
#Create test and train data sets
```

```
trainSet <- subCredit[trainList,]
```

```
testSet <- subCredit[-trainList,]
```


Train Using a Treebag Model

```
fit1 <- train(Class ~ ., data=trainSet, method="treebag",  
              preProc=c("center","scale"))
```

The **train()** command trains the specified model (here it is “**treebag**”)

Class ~ . This is the standard model specification syntax; the dot after the tilde includes all of the other variables as predictors; otherwise spell them out separated with plus signs

preProc= allows preprocessing of the input data, in this case taking the precaution of centering and scaling to put every input variable on the same scale

Interpret Results: Variable Importance

```
varImp(fit1)
```

```
treebag variable importance
```

	Overall
Amount	100.00
Age	76.73
Duration	57.75
ResidenceDuration	35.73
InstallmentRatePercentage	34.33
NumberExistingCredits	20.84
Telephone	13.91
NumberPeopleMaintenance	11.77
ForeignWorker	0.00

Interpret Results: Use rpart

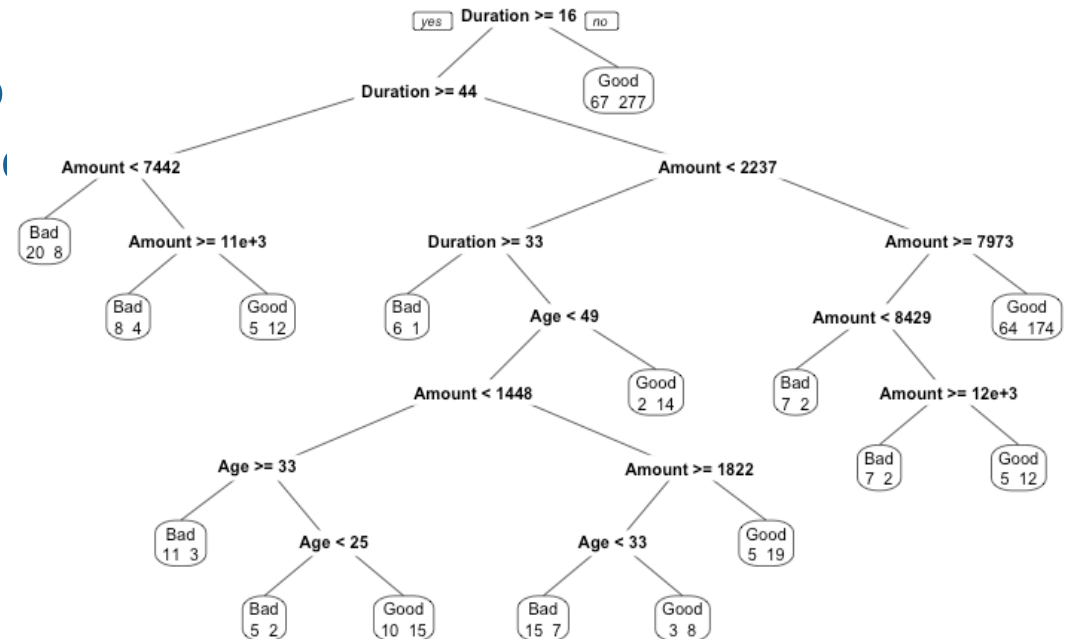
Use rpart to display a simple CART tree with key variables

This tree gives a general idea of how variables are working, but is different from the treebag model (which has 25 trees!)

```
cartTree <- rpart(Class ~ Amount + Age + Duration  
  data = trainSet, method="
```

```
prp(cartTree, faclen = 0, cex = 0.8, extra = 1)
```

treebag variable importance	
	Overall
Amount	100.00
Age	76.73
Duration	57.75
. . .	



Step 4: Assess Fit With New (Test) Data

```
predOut <- predict(fit1, newdata=testSet)
confusion <- confusionMatrix(predOut, testSet$Class)
confusion
```

Confusion matrix and statistics

	Reference	
Prediction	Bad	Good
Bad	20	18
Good	40	122

Accuracy : 0.71

Question

Is this a good model?

Confusion matrix and statistics

		Reference	
Prediction	Bad	Good	
	Bad	20	18
Good	40	122	

Accuracy : 0.71

95% CI : (0.6418, 0.7718)

No information rate : 0.7

p value [Acc > NIR] : 0.412322



Advanced Tree Example (cont.)

School of Information Studies
Syracuse University

Interpret Results: Use rpart

No Information Rate: 0.7

Model Accuracy: 0.71

p value [Acc > NIR] : 0.412322

→ this suggests the improvement is not significant




Using Support Vector Machines

School of Information Studies
Syracuse University


Visualizing the Algorithm, Part I

This is easy



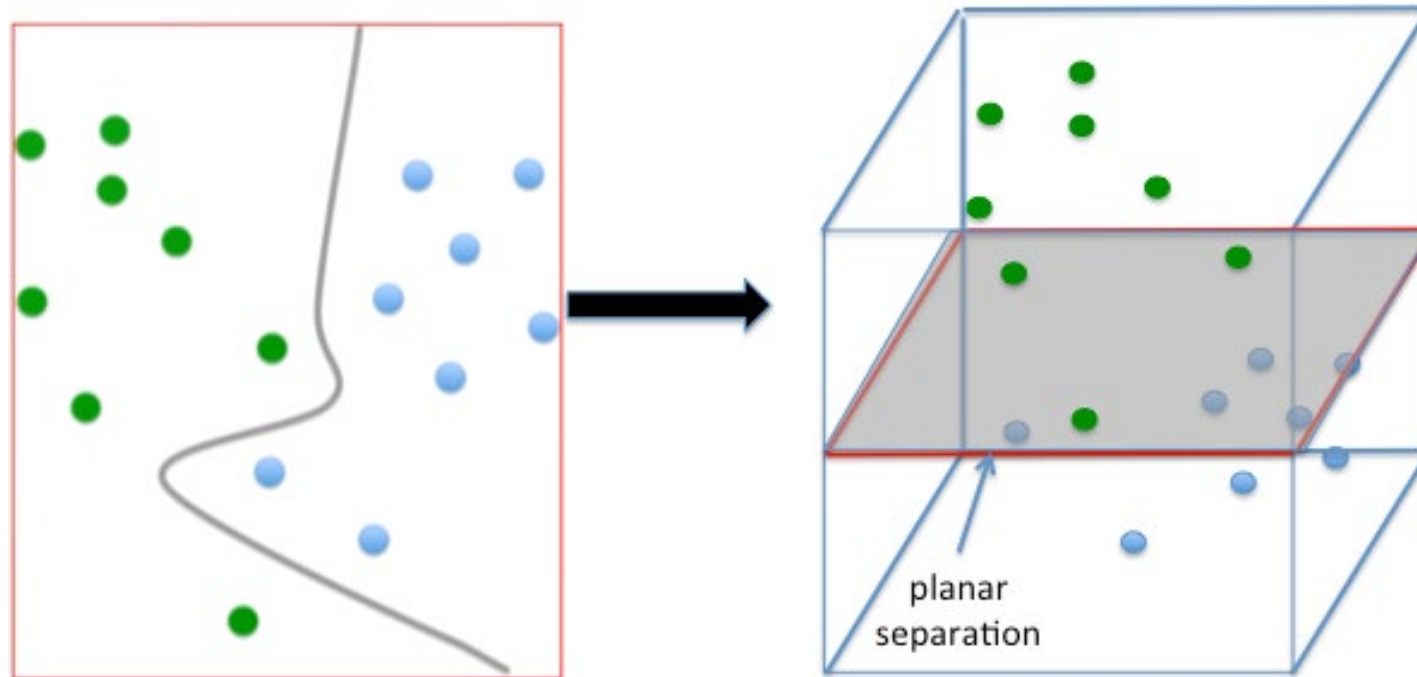
Visualizing the Algorithm, Part II

What about this?



Visualizing the Algorithm, Part III

2D to 3D mapping




SVM: What Are Support Vectors?

Very **different approach to prediction** compared with the linear methods and tree methods

Based on **statistical learning theory**

Use the **cases that are difficult to predict** to help define a robust prediction model (because any model can handle “easy” cases!)

Support vectors **are harder-to-predict cases** (on a defined border and margin that separate the cases)



Use SVM to Explore the Credit Data

#Step 1: load the required packages that were already done

#Step 2: Create training and test sets

```
data("GermanCredit")
```

```
subCredit <- GermanCredit[,1:10]
```

#Makes the sampling predictable

```
set.seed(111)
```

#Randomly sample elements to go into a training data set

```
trainList <- createDataPartition(y=subCredit$Class,p=.80,list=FALSE)
```

```
trainSet <- subCredit[trainList,]
```

```
testSet <- subCredit[-trainList,]
```

Use SVM to Explore the Credit Data (cont.)

#Train the SVM model

```
fit2 <- train(Class ~ ., data=trainSet, method="svmRadial",  
              preProc=c("center","scale"))
```

#Assess fit with new data

```
predOut <- predict(fit2, newdata=testSet)
```

Compare the Results

TreeBag

```
confMatrix <- table(predOut, testSet$Class)
confMatrix

predOut  Bad  Good
   Bad    20   18
   Good   40  122

prop.table(confMatrix)

predOut  Bad  Good
   Bad  0.10 0.09
   Good 0.20 0.61

errorRate <- (sum(confMatrix) -
              sum(diag(confMatrix))) /
              sum(confMatrix)

errorRate
[1] 0.29
```

SVM

```
confMatrix <- table(predOut, testSet$Class)
confMatrix

predOut  Bad  Good
   Bad    10    3
   Good   50  137

prop.table(confMatrix)

predOut  Bad  Good
   Bad  0.05 0.015
   Good 0.25 0.685

errorRate <- (sum(confMatrix) -
              sum(diag(confMatrix))) /
              sum(confMatrix)

errorRate
[1] 0.265
```

Compare the Results (Varimp)

TreeBag

	Importance
Amount	100.00
Age	76.73
Duration	57.75
ResidenceDuration	35.73
InstallmentRatePerc	34.33
NumberExistingCredits	20.84
Telephone	13.91
NumberPeopleMaint.	11.77
ForeignWorker	0.00

SVM

	Importance
Duration	100.00
Age	51.92
Amount	43.20
InstallmentRatePerc	31.54
NumberExistingCredits	25.19
ForeignWorker	9.85
NumberPeopleMaint	7.01
ResidenceDuration	0.27
Telephone	0.00

Explore the Confusion Matrix

#Review the error—use the built in 'confusionMatrix'

```
confusion <- confusionMatrix(predOut, testSet$Class)
```

```
confusion
```

Confusion matrix and statistics

	Reference	
Prediction	Bad	Good
Bad	10	3
Good	50	137

Accuracy : 0.735

95% CI : (0.6681, 0.7948)

No information rate : 0.7

p value [Acc > NIR] : 0.1579

Explore the “C” Parameter

fit2

Support vector machines with radial basis function kernel

800 samples

9 predictor

2 classes: 'Bad', 'Good'

Preprocessing: centered (9), scaled (9)


...

C	Accuracy
0.25	0.6969592
0.50	0.6974949
1.00	0.6989207

Why Error Might Not Be Bad

Non-separable training sets

Use linear separation, but admit training errors.



Penalty of error: distance to hyperplane multiplied by *error cost* C .

Cost Parameter ("C")

KSVM cost parameter impact summary

Higher cost "C" value	Fewer classification mistakes Fewer problem points	Smaller margin of separation	Specialized model	Higher cross-validation error	Lower training error
Lower cost "C" value	More classification mistakes More problem points	Higher margin of separation	Generalized model	Lower cross-validation error	Higher training error

Question

If SVM is such a “black box”—where it is hard to know how the model actually works and what variables matter the most—why would we ever bother to use it?



Using Support Vector Machines (cont.)

School of Information Studies
Syracuse University

Answer

If SVM is such a "black box"—where it is hard to know how the model actually works and what variables matter the most—why would we ever bother to use it?

Sometimes just need the best model possible

Sometimes can try to explain the model by showing how the model works with some specific examples