



Introduction to Text Mining

School of Information Studies
Syracuse University

An Example of Unstructured Data

Friends and fellow citizens: I stand before you tonight under indictment for the alleged crime of having voted at the last presidential election, without having a lawful right to vote. It shall be my work this evening to prove to you that in thus voting, I not only committed no crime, but, instead, simply exercised my citizen's rights, guaranteed to me and all United States citizens by the National Constitution, beyond the power of any state to deny.

The preamble of the Federal Constitution says:

- “We, the people of the United States, in order to form a more perfect union, establish justice, insure domestic tranquility, provide for the common defense, promote the general welfare, and secure the blessings of liberty to ourselves and our posterity, do ordain and establish this Constitution for the United States of America.”
- It was we, the people; not we, the white male citizens; nor yet we, the male citizens; but we, the whole people, who formed the Union. And we formed it, not to give the blessings of liberty, but to secure them; not to the half of ourselves and the half of our posterity, but to the whole people—women as well as men. And it is a downright mockery to talk to women of their enjoyment of the blessings of liberty while they are denied the use of the only means of securing them provided by this democratic-republican government—the ballot.

Inside the CPU, Text Is Easy

From a computer perspective—binary encoding for storing written language symbols from the world’s various systems

Binary	Oct	Dec	Hex	Glyph		
				1963	1965	1967
010 0000	040	32	20			space
010 0001	041	33	21			!
010 0010	042	34	22			"
010 0011	043	35	23			#
010 0100	044	36	24			\$
010 0101	045	37	25			%
010 0110	046	38	26			&
010 0111	047	39	27			'
010 1000	050	40	28			(
010 1001	051	41	29)
010 1010	052	42	2A			*
010 1011	053	43	2B			+
010 1100	054	44	2C			,
010 1101	055	45	2D			-
010 1110	056	46	2E			.
010 1111	057	47	2F			/
011 0000	060	48	30			0
011 0001	061	49	31			1
011 0010	062	50	32			2

Row	Cells	Range(s)
00	20–7E	Basic Latin (00–7F)
	A0–FF	Latin-1 Supplement (80–FF)
01	00–13, 14–15, 16–2B, 2C–2D, 2E–4D, 4E–4F, 50–7E, 7F	Latin Extended-A (00–7F)
	8F, 92, B7, DE–EF, FA–FF	Latin Extended-B (80–FF ...)
02	18–1B, 1E–1F	Latin Extended-B (... 00–4F)
	59, 7C, 92	IPA Extensions (50–AF)
03	BB–BD, C6, C7, C9, D6, D8–DB, DC, DD, DF, EE	Spacing Modifier Letters (B0–FF)
	74–75, 7A, 7E, 84–8A, 8C, 8E–A1, A3–CE, D7, DA–E1	Greek (70–FF)
04	00–5F, 90–91, 92–C4, C7–C8, CB–CC, D0–EB, EE–F5, F8–F9	Cyrillic (00–FF)
1E	02–03, 0A–0B, 1E–1F, 40–41, 56–57, 60–61, 6A–6B, 80–85, 9B, F2–F3	Latin Extended Additional (00–FF)
1F	00–15, 18–1D, 20–45, 48–4D, 50–57, 59, 5B, 5D, 5F–7D, 80–B4, B6–C4, C6–D3, D6–DB, DD–EF, F2–F4, F6–FE	Greek Extended (00–FF)
20	13–14, 15, 17, 18–19, 1A–1B, 1C–1D, 1E, 20–22, 26, 30, 32–33, 39–3A, 3C, 3E, 44, 4A	General Punctuation (00–6F)
	7F, 82	Superscripts and Subscripts (70–9F)
	A3–A4, A7, AC, AF	Currency Symbols (A0–CF)

From ASCII (7-bits) to Unicode: An international standard that supports up to four bytes (32 bits) per symbol and that encodes 137,439 characters from 146 different scripts

(Source: Wikipedia)

Outside the CPU, Text Is Difficult

From an analyst's perspective:

- Documents are organized into a corpus
- Each document contains an encoded sample of written language, generally:
 - Human readable files (e.g., plain text)
 - Metadata describing: Data source, date of capture, speaker/author, etc.
 - Character/symbol encoding appropriate to language
 - One file per document, or container file includes document separator characters (e.g., newline)
- Each file contains unstructured, “natural” language content—generally the same “type” of content for each file in the corpus (i.e., don’t mix tweets with books)



Text Analytics Overview

Text mining (TM):

- Looking for unexpected patterns in large data sets
- Focuses on statistical approaches such as counting word frequencies

Natural language processing (NLP):

- Understand how to program machines to make sense of human language
- Use linguistics to break text into its grammatical pieces, such as nouns and verbs

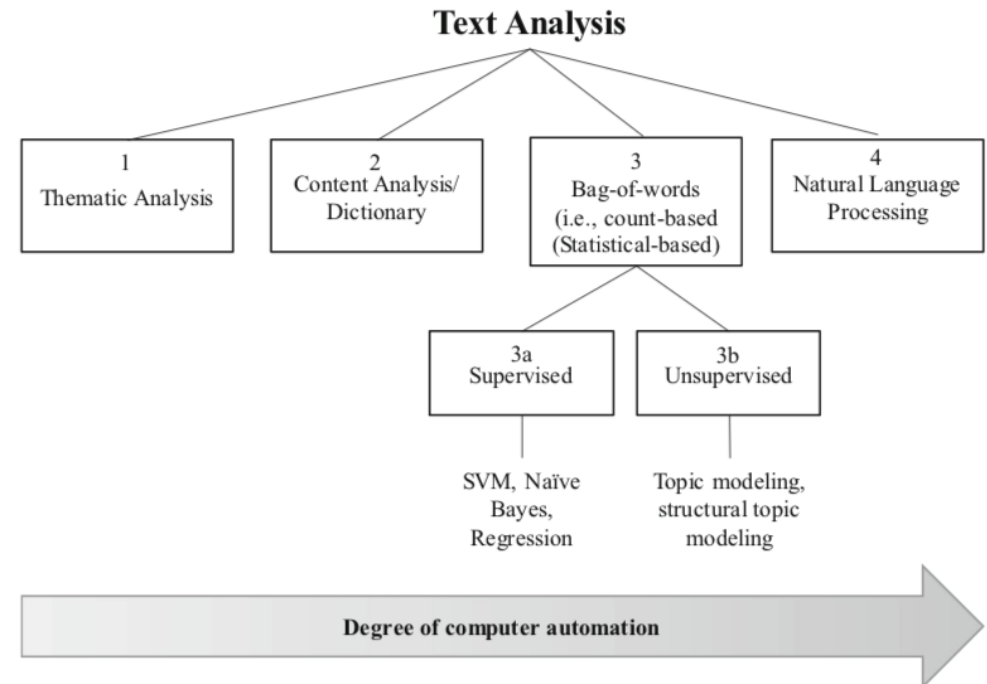


Image credit: Banks et al. (2018), Fig. 1, p. 447

“Unstructured” Is the Key

Natural language is notoriously flexible and ambiguous—even simple tasks like parsing are complicated.

- Contractions; compound words; misspellings; proper names; preposition attachment; vernacular

Transforming a set of documents into data that allow for comparisons, linkages, visualization, or other analysis, is **complex and contains many decisions**.




Image credit: [NLTK.org](https://nltk.org)

Example Challenge in Text Analysis: Preposition Attachment

I saw the man on the hill with a telescope.



1. I saw the man. The man was on the hill. I was using a telescope.
2. I saw the man. I was on the hill. I was using a telescope.
3. I saw the man. The man was on the hill. The hill had a telescope.
4. I saw the man. I was on the hill. The hill had a telescope.
5. I saw the man. The man was on the hill. I saw him using a telescope.

Image credit: Zareen Syed, UMBC

Question

If automated text analysis is so tricky, why bother? Why not just measure clicks and do surveys?



Introduction to Text Mining (cont.)

School of Information Studies
Syracuse University

Question

If automated text analysis is so tricky, why bother? Why not just measure clicks and do surveys?



The Term-Document Matrix (AKA Document-Term Matrix)

School of Information Studies
Syracuse University

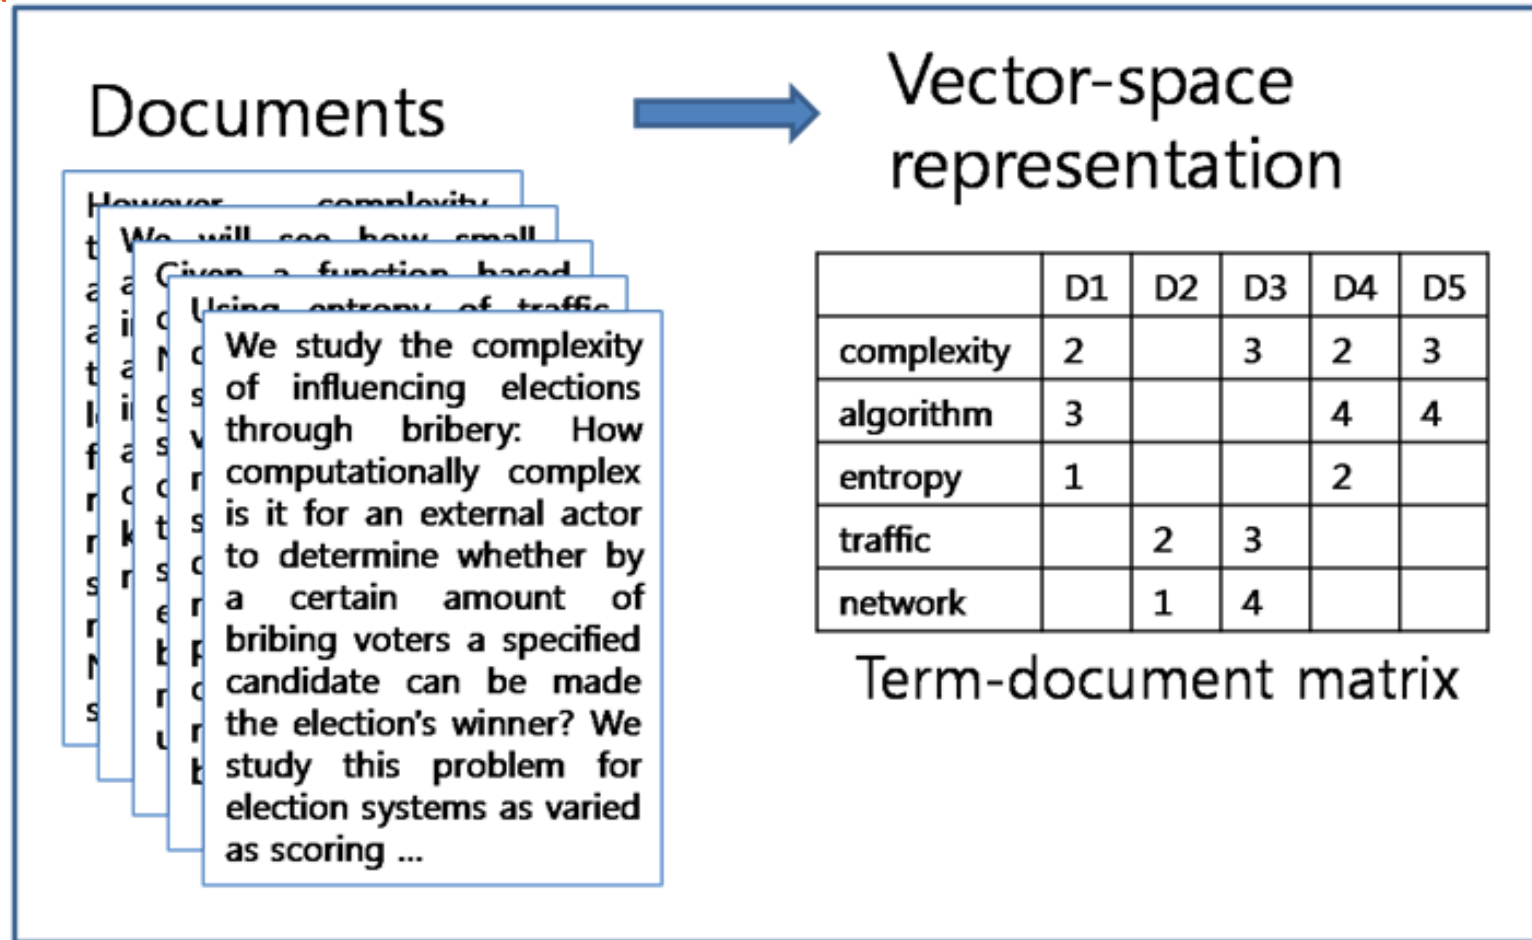
Statistical Approaches to Text

	I	love	dogs	hate	and	knitting	is	my	hobby	passion
Doc 1	1	1	1							
Doc 2	1		1	1	1	1				
Doc 3					1	1	1	2	1	1

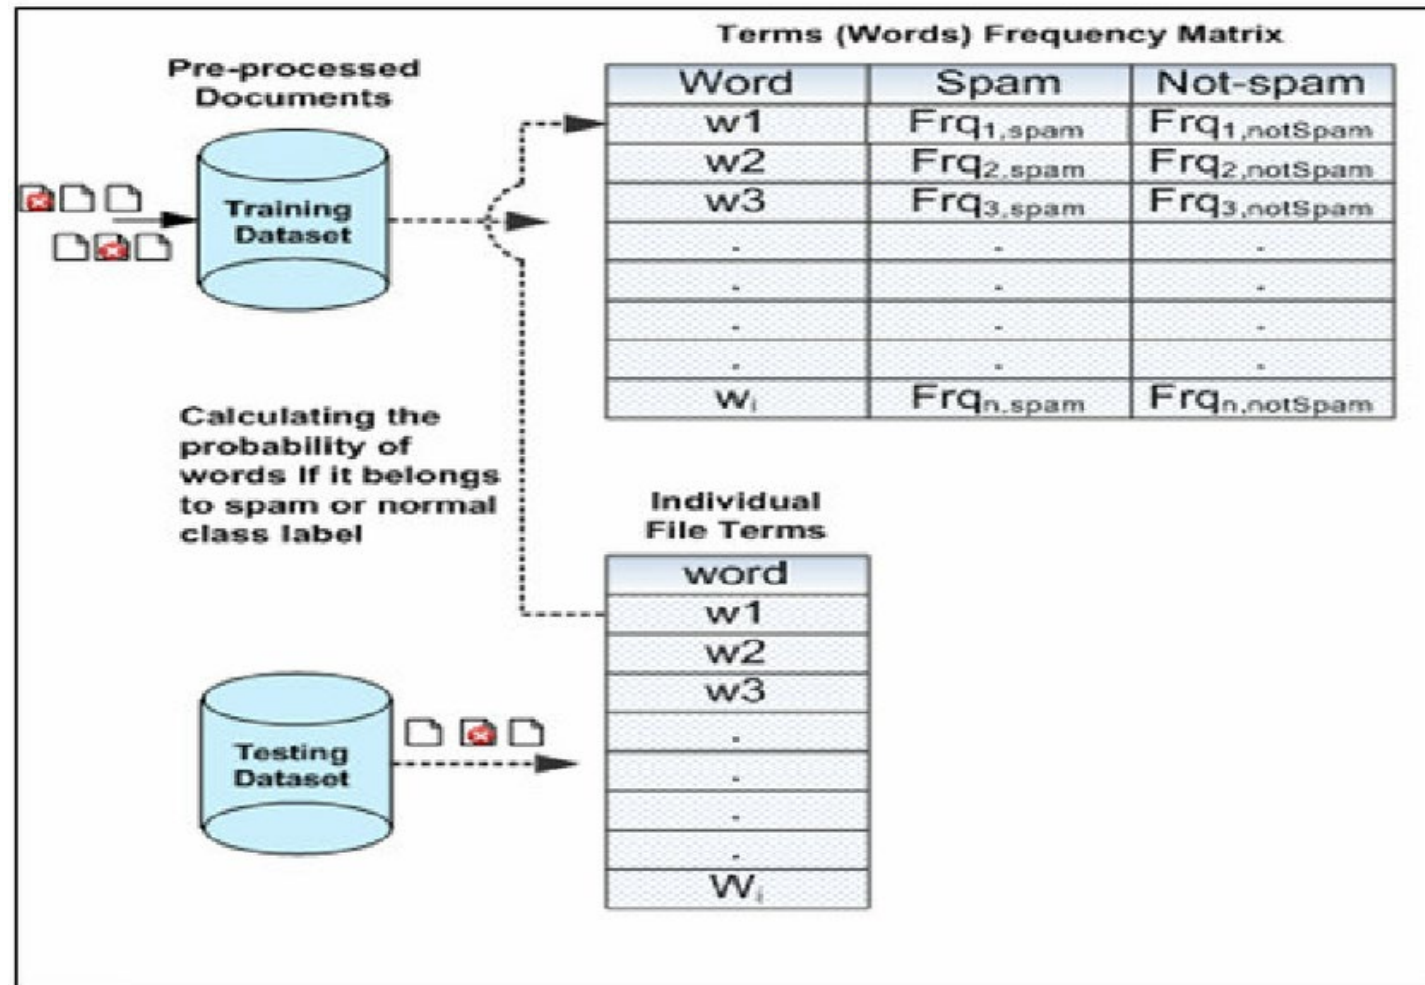
Image credit: Pio Calrderon

- Treat each text fragment as a collection of units/words, without regard to word order
- Use many text fragments – as small as a phrase or as large as a book; each fragment is considered a “document”
- Compute a term-document matrix (TDM) or document-term matrix; a sparse matrix flagging the appearance of a term in a document
- Conduct statistical or machine learning analysis on the TDM to reveal patterns

Example TDM



Example Application: Spam Filter



Question

Describe a useful application of a term-document matrix?



Creating Word Clouds

School of Information Studies
Syracuse University

Read Some Web Text

```
url <- "http://programdevelopment.syr.edu/wp-  
content/uploads/2018/03/REM_Manuscript_fall2017_text.txt"  
charVector <- scan(url, character(0), sep = "\n")  
head(charVector)
```

```
'SYRACUSE MANUSCRIPT'  
'FALL 2017 | VOL. 6 | NO. 2'  
'SYRACUSE UNIVERSITY\'S AFRICAN AMERICAN AND LATINO ALUMNI MAGAZINE'  
'ON THE COVER:'  
'Left to right, from top: Cheryl Wills \'89 and Taye Diggs \'93; Lazarus Sims  
'\'96; Lt. Col. Pia W. Rogers \'98, G\'01, L\'01 and Dr. Akima H. Rogers \'94;'  
'Amber Hunter \'19, Nerys Castillo-Santana \'19, and Nordia Mullings \'19; Demaris  
'Mercado \'92; Dr. Ruth Chen and Chancellor Kent Syverud; Carmelo Anthony; Darlene  
'Harris \'84 and Debbie Harris \'84 with Soledad O\'Brien'
```

Library(quanteda)

The quanteda package was developed by Kenneth Benoit ([Kenbenoit](#)) at the London School of Economics.

Does everything that TM does (an older package), plus much more:

- Includes readtext(), a wrapper for a wide range of text files (PDF, CSV, text, XML, JSON) that simplifies reading in folders of text files
- Directly supports document-level variables for machine learning
- A comprehensive text processing package that uses some C++ and Fortran—but unlike Stanford Core NLP tools, no Java; mostly runs faster and with less memory than Stanford Core NLP

Text Transformations

Stemming (removing the endings of words)

Removing the punctuation

Removing low frequency terms

Taking out the “stop” words

Words such as “*the*,” “*a*,” and “*at*” appear in so many different parts of the text that they are useless for differentiating between documents.

Additional possible transformations not shown here: lowercase, removing numbers

Preparing and Visualizing Corpus

#Make a corpus of web data

```
intcorpus <- corpus(charVector)
```

```
paras <- corpus_reshape(intcorpus, to="paragraphs")
```

#Turn the corpus into a document-term matrix while removing stop words and punctuation

```
webfile_dtm <- dfm(paras, stem=TRUE, remove_punct=TRUE, remove=stopwords("english"))
```

```
webfile_dtm <- dfm_trim(webfile_dtm, min_termfreq=2)
```


#2,079 documents, 976 features (99.6% sparse)

```
webfile_dtm
```

#Visualize words

```
textplot_wordcloud(webfile_dtm)
```

```
textplot_wordcloud(webfile_dtm)
```



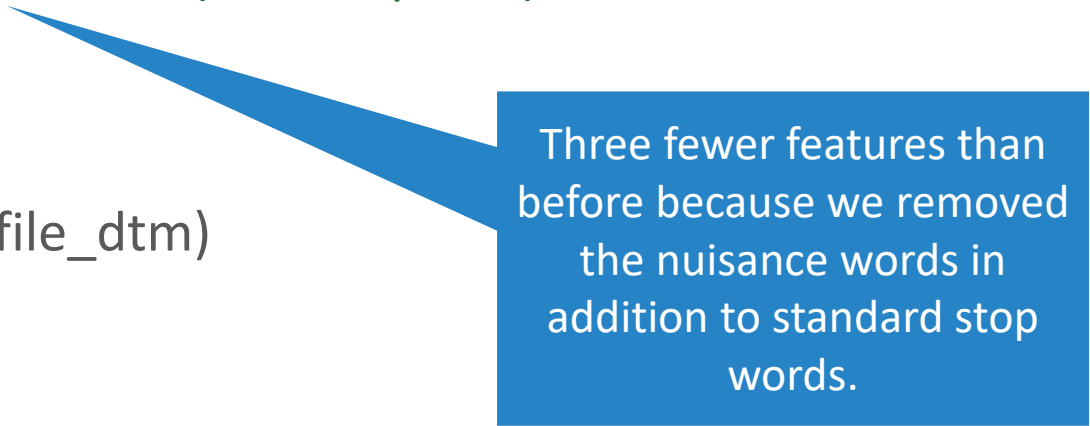
Let's Trim Out Some Useless Words

```
webfile_dtm <- dfm(paras, stem=TRUE, remove_punct=TRUE,  
                  remove=c(stopwords("english"), "ms", "mr", "mrs"))  
webfile_dtm <- dfm_trim(webfile_dtm, min_termfreq=3)
```

#2,079 documents, 973 features (99.6% sparse)

webfile_dtm

textplot_wordcloud(webfile_dtm)



Three fewer features than before because we removed the nuisance words in addition to standard stop words.

```
textplot_wordcloud(webfile_dtm)
```



Let's Raise the Bar

#Raise the frequency threshold and turn off stemming

```
webfile_dtm <- dfm(paras, stem=FALSE, remove_punct=TRUE,  
                  remove=c(stopwords("english"), "ms", "mr", "mrs"))
```

#Set a minimum word frequency

```
webfile_dtm <- dfm_trim(webfile_dtm, min_termfreq=18)
```

```
textplot_wordcloud(webfile_dtm)
```

By raising min_termfreq, we eliminate words from the TDM that are less frequent. This can be dangerous for indexing! Infrequent words can be powerful!



Question

Justify when it might be appropriate to use a word cloud as part of a text analysis.



Creating Word Clouds (cont.)

School of Information Studies
Syracuse University

Answer

Justify when it might be appropriate to use a word cloud as part of a text analysis.

- Example: get a “feeling” for answers from a survey response



Sentiment Analysis

School of Information Studies
Syracuse University

Conceptual Methodology

Load positive and negative word lists

Count positive words and negative words that match those in the word list from the TDM

Compute ratios: positive/total and negative/total

```
#Reprocess the TDM to include low frequency words
webfile_dtm <- dfm(paras, stem=FALSE, remove_punct=TRUE,
                  remove=c(stopwords("english"), "ms", "mr", "mrs"))
webfile_dtm <- dfm_trim(webfile_dtm, min_termfreq=1)

#2,079 documents, 3,735 features (99.9% sparse)
webfile_dtm
```

1. Make a “Named Number” List of Words

#Create a simple matrix from the TDM

```
webWords <- as.matrix(webfile_dtm)  
str(webWords)
```

All we have done is simplify the more complex—and more efficiently stored TDM—into a regular rectangular matrix.

```
num [1:2079, 1:3735] 1 0 1 0 0 0 0 0 0 0 ...  
- attr(*, "dimnames")=List of 2  
..$ docs      : chr [1:2079] "text1.1" "text2.1" "text3.1" "text4.1" ...  
..$ features: chr [1:3735] "syracuse" "manuscript" "fall" "2017" ...
```

Note that the columns of this matrix are labeled with the terms from the TDM.

1. Make a “Named Number” List of Words (cont.)

#Sum the columns to get the overall term counts

```
wordCounts <- colSums(webWords)
```

```
wordCounts <- sort(wordCounts, decreasing=TRUE)
```

#Check the first several items in "wordCounts"

```
head(wordCounts, 20)
```

says	scholarship	alumni	syracuse	cbt	students
107	104	100	85	77	76
g	su	university	endowed	1	fund
70	58	57	56	51	50
m	dr	black	program	miss	89
47	44	43	42	42	37

These are “named” numbers: a numeric vector but with a text attribute that stores a label for each value

2. Read the Positive and Negative Dictionaries

#2006 items

```
posWords <- scan("positive-words.txt", character(0), sep = "\n")
```

#4783 items

```
negWords <- scan("negative-words.txt", character(0), sep = "\n")
```

```
head(posWords, 18)
```

[1]	"a"	"abound"	"abounds"	"abundance"
[5]	"abundant"	"accessible"	"accessible"	"acclaim"
[9]	"acclaimed"	"acclamation"	"accolade"	"accolades"
[13]	"accommodative"	"accomodative"	"accomplish"	"accomplished"
[17]	"accomplishment"	"accomplishments"		

3. Match Positive and Negative Words

```
matchedP <- match(names(wordCounts), posWords, nomatch = 0)
```

```
matchedN <- match(names(wordCounts), negWords, nomatch = 0)
```

- `nomatch=0` specifies what value ("0") to use when an item has no match in the second vector

Here are the first 100 values of

`matchedP`:

```
head(matchedP, 100)
```

[1]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[16]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[31]	0	0	0	0	0	0	0	0	0	0	1746	0	0	0	0
[46]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[61]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[76]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[91]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

```
Names (wordCounts) [41]
```

```
[1] "support"
```

```
posWords [1746]
```

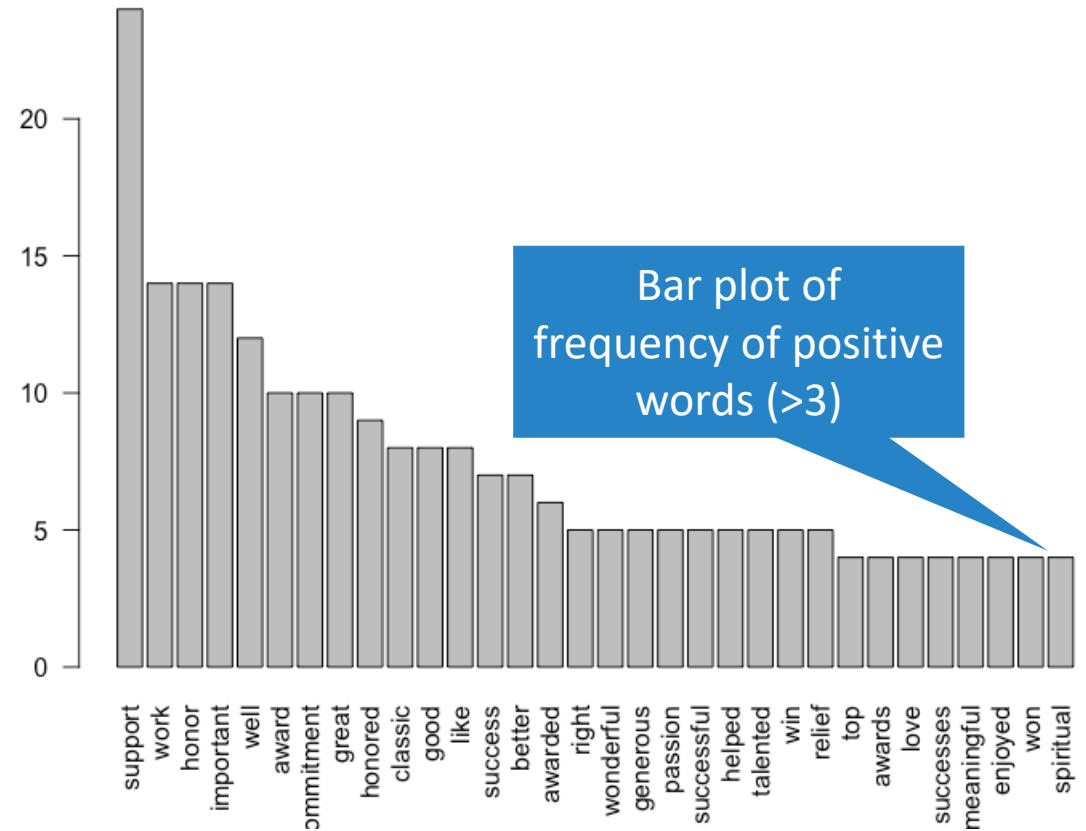
```
[1] "support"
```

Item 41 in `names(wordCounts)` matches with item 1746 in `posWords`.

Matched?

matchedP is mostly empty space—lots of zeros.

Where matchedP is not empty, it just contains an index reference to posWords.



4. Calculate the Proportions!

#Calculate the proportion of words that are positive or negative

```
sum(matchedP)/totalWords
```

```
[1] 0.04228368
```

```
sum(matchedN)/totalWords
```

```
[1] 0.01248885
```

This document contains about 3.4 times as many positive words as negative words.

- Is this a surprising result? Why or why not?

Questions

Does this truly measure sentiment? Where could it go wrong?

What are some circumstances where we might draw the wrong conclusions?



Sentiment Analysis (cont.)

School of Information Studies
Syracuse University

Answers

Does this truly measure sentiment? Where could it go wrong?

What are some circumstances where we might draw the wrong conclusions?

How about:

- “No, that movie was not bad.”
- “I liked the movie. The lead was a terrible person, and his partner was even worse.”